

Understanding Document References



Dan Geabunea

SENIOR SOFTWARE DEVELOPER

@romaniancoder www.romaniancoder.com



Overview



Mongo document references

Creating document references in Spring
using @DbRef

Spring Mongo lifecycle events

Demo: Implementing document
references & cascading



Mongo Document References



Document Reference

A way to link together multiple documents that are related



For most cases, the denormalized model where data is stored in a single document is optimal



A Simple Example

Bank Transaction 1

```
{  
  "_id": 1,  
  "amount": 3000,  
  "customer": {  
    "name": "John Doe",  
    "nationalIdNb": "AZ456734"  
  }  
}
```

Bank Transaction 2

```
{  
  "_id": 2,  
  "amount": 1000,  
  "customer": {  
    "name": "John Doe",  
    "nationalIdNb": "AZ456734"  
  }  
}
```

The Problem

Bank Transaction 1

```
{
  "_id": 1,
  "amount": 3000,
  "customer": {
    "name": "John Doe",
    "nationalIdNb": "BG90865"
  }
}
```

Bank Transaction 2

```
{
  "_id": 2,
  "amount": 1000,
  "customer": {
    "name": "John Doe",
    "nationalIdNb": "AZ456734"
  }
}
```

Relationships

Customer

```
{  
  "_id": 234,  
  "name": "John Doe",  
  "nationalIdNb": "BG90865"  
}
```

Bank Transaction

```
{  
  "_id": 849,  
  "amount": 1000,  
  "customerId": 234  
}
```


Relating Documents

Manual

Saving the `'_id'` of a document in another document

DBRefs

Link documents together by using the `'_id'` field, collection name and database name.




DBRef



Can link documents across collections or across different databases



To resolve DBRefs, the application must perform additional queries



Cascading is not supported



DBRef Format

\$ref

The name of the collection where the linked document resides

\$id

The value of the `_id` field of the referenced document

\$db

The name of the database where the referenced document resides



Aircraft

```
{  
  "_id": ObjectId("4e08277f-2660-43b6-ae71-ef2f5a4c9ce2")  
  "name": 737  
  "capacity": 234  
  "engine" : {  
    "$ref": "engines",  
    "$id": ObjectId("3e3df0ff-5531-4e9c-8a90-10e2a6faef8a"),  
    "$db": "atm"    // optional  
  }  
}
```

The order of fields in the DBRef matters and must be used in the correct sequence



DBRefs are not “smart”



Creating Document References in Spring Using @DBRef



Aircraft

@Document

```
public class Aircraft {  
    @Id private String id;  
    private String model;  
    private int maxPassengers;  
    private Engine engine; // embedded document  
}
```


Aircraft

@Document

```
public class Aircraft {  
    @Id private String id;  
    private String model;  
    private int maxPassengers;  
    @DBRef private Engine engine;    // document reference  
}
```

Engine

@Document

```
public class Engine {  
    @Id private String id;  
    private String type;  
    private int maxPower;  
}
```

Spring Data MongoDB
fetches documents
annotated with @DBRef for
us automatically



```
Aircraft a = new Aircraft("737");  
a.setEngine(new Engine("GM", 1000));  
  
aircraftRepository.insert(a);
```

Insert

You won't get an error

But the engine is not getting saved in the Engine collection



Cascading does not work
with DBRefs on save by
default



```
Engine e = new Engine("GM", 1000);  
engineRepository.insert(e);  
  
Aircraft a = new Aircraft("737");  
a.setEngine(e);  
aircraftRepository.insert(a);
```

The Correct Way

Save the (child) object first

Save the (parent/root) after that



```
@Document
public class Aircraft {
    @Id private String id;
    @DBRef(lazy = true) private Engine engine;
}
```

Lazy Loading

By default, `@DBRef` instructs the framework to perform an eager load

This can be changed by modifying the `lazy` property to “true”


- The sub-document is loaded when one of its fields/methods is accessed
- Useful for bi-directional references to avoid a cyclic dependency



@DBRef



@DBRef to creates document references



Eager loading by default but can be changed. The Spring framework fetches the referenced document for you and populates the POJO



Cascading on save is not available



You can hook into Mongo lifecycle events and implement cascading



Pay attention to queries on
related documents



Filtering

Aircraft

@Document

```
public class Aircraft {  
    @Id private String id;  
    private Engine engine;  
}
```

Repository

```
@Query("{ 'engine.maxPower': ?0 }")  
List<Aircraft> byPower(int pow);  
  
// Will work
```

Filtering

Aircraft

```
@Document  
public class Aircraft {  
    @Id private String id;  
    @DBRef private Engine engine;  
}
```

Repository

```
@Query("{ 'engine.maxPower': ?0 }")  
List<Aircraft> byPower(int pow);  
  
// Won't work
```

Spring Mongo Lifecycle Events



Lifecycle Events

Events that your application can respond to by registering special beans in the Spring application context.



“Before” Events

onBeforeConvert

onBeforeSave

onBeforeDelete



“After” Events

onAfterConvert

onAfterSave

onAfterLoad

onAfterDelete



Save / Update Document Events



Call insert or save in MongoTemplate



onBeforeConvert is called before the Java object is converted to a Document by the MongoConverter



onBeforeSave is called before inserting or saving the document in the database



onAfterSave is called after the document was inserted or saved in the database



Load Document Events



Call `find`, `findOne`, `findAndRemove` in `MongoTemplate`



`onAfterLoad` is called after the Document has been retrieved from the database



`onAfterConvert` is called after the Document has been converted into a POJO



Delete Document Events



Call `remove` on `MongoTemplate`



`onBeforeDelete` is called just before the document gets deleted from the database



`onAfterDelete` is called after the Document has been deleted from the database



Event Behavior

Lifecycle events are emitted only for root level types. Sub-documents are not subject to event publication unless they are annotated with `@DBRef`

It is all happening in an async fashion. We have no guarantees to when an event is processed.



“Why would I ever want to hook into these events?”


You



Use Cases



Implement cascade on
save



Trigger some
job/action in different
systems



Security audits



“All right, so how can I hook into these events and execute my custom logic”

You



AbstractMongoEventListener

```
public abstract class AbstractMongoEventListener<E>
    implements ApplicationListener<MongoMappingEvent<?>> {

    public void onAfterConvert(AfterConvertEvent<E> event) {...}
    public void onAfterDelete(AfterDeleteEvent<E> event) {...}
    public void onAfterLoad(AfterLoadEvent<E> event) {...}

    // ...

}
```

BeforeSaveListener

```
public class BeforeConvertListener
    extends AbstractMongoEventListener<Flight> {

    @Override
    public void onBeforeConvert(BeforeConvertEvent<Flight> event) {
        // your custom logic...
    }
}
```


BeforeSaveListener

@Component

```
public class BeforeSaveListener
    extends AbstractMongoEventListener<Flight> {

    @Override
    public void onBeforeConvert(BeforeConvertEvent<Flight> event) {
        // your custom logic...
    }
}
```

Create an event listener per
feature



Demo



Demo: Implementing Document References

- @DBRef
- Implement cascading and auditing using lifecycle events



Summary



Documents that are related can be linked via manual references or by DBRefs

The Spring Data Mongo framework does things under the hood to make your experience easier with @DBRef

Cascading on save is not implemented for DBRefs, but you can achieve this by hooking into lifecycle events

Mongo lifecycle events are a great way to add valuable logic in key persistence moments



For nearly every case where you want to store a relationship between two documents, **use manual references**. Your application can resolve references as needed.

MongoDB official documentation



However, if you need to reference documents from **multiple collections/databases**, consider using DBRefs.

MongoDB official documentation



From My Experience Using Mongo with Spring

Start by designing data structures without relationships. Stay true to the NoSQL philosophy.

Use @DBRef but only if you implement cascade on save and you don't create bidirectional references

- Population of related documents done by the framework
- Cascading

Manual references

